

Introduction to deep learning

Lecture 3

Sam van Gool
vangool@irif.fr

Master 2 informatique
Université Paris Cité

21 January 2025

Today's lecture

- 1 Recap
- 2 Softmax regression and SGD: implementation
- 3 Vectorization
- 4 Going deep: Activation functions

Last week

CM:

- Softmax regression

TP:

Reminder: Course webpage: <https://samvangool.net/iap/>

Last week

CM:

- Softmax regression
- Intuition for *cross entropy loss*

TP:

Reminder: Course webpage: <https://samvangool.net/iap/>

Last week

CM:

- Softmax regression
- Intuition for *cross entropy loss*
- Different forms of *gradient descent*:

TP:

Reminder: Course webpage: <https://samvangool.net/iap/>

Last week

CM:

- Softmax regression
- Intuition for *cross entropy loss*
- Different forms of *gradient descent*:
 - naive (pass through all the data before making an update)

TP:

Reminder: Course webpage: <https://samvangool.net/iap/>

Last week

CM:

- Softmax regression
- Intuition for *cross entropy loss*
- Different forms of *gradient descent*:
 - naive (pass through all the data before making an update)
 - stochastic minibatch.

TP:

Reminder: Course webpage: <https://samvangool.net/iap/>

Last week

CM:

- Softmax regression
- Intuition for *cross entropy loss*
- Different forms of *gradient descent*:
 - naive (pass through all the data before making an update)
 - stochastic minibatch.

TP:

- Implementing *softmax regression* in Numpy

Reminder: Course webpage: <https://samvangool.net/iap/>

Last week

CM:

- Softmax regression
- Intuition for *cross entropy loss*
- Different forms of *gradient descent*:
 - naive (pass through all the data before making an update)
 - stochastic minibatch.

TP:

- Implementing *softmax regression* in Numpy
- Over-fitting a single batch

Reminder: Course webpage: <https://samvangool.net/iap/>

Last week

CM:

- Softmax regression
- Intuition for *cross entropy loss*
- Different forms of *gradient descent*:
 - naive (pass through all the data before making an update)
 - stochastic minibatch.

TP:

- Implementing *softmax regression* in Numpy
- Over-fitting a single batch
- Training on a full dataset, achieving $> 80\%$ accuracy.

Reminder: Course webpage: <https://samvangool.net/iap/>

TO DO last week

- Finish TP2.

TO DO last week

- Finish TP2.
- Read [Z, pp. 82-88, pp. 126-148], [P, Sections 5.1–5.2].

TO DO last week

- Finish TP2.
- Read [Z, pp. 82-88, pp. 126-148], [P, Sections 5.1–5.2].
- Do [Z, Section 4.4.7], Exercise 1 and [P, Problem 5.2].

TO DO last week

- Finish TP2.
- Read [Z, pp. 82-88, pp. 126-148], [P, Sections 5.1–5.2].
- Do [Z, Section 4.4.7], Exercise 1 and [P, Problem 5.2].
- If you need a reminder of Linear Algebra and Calculus:
read [Z, Section 2.3 and 2.4], or [P, Appendix B].

Today's lecture

- 1 Recap
- 2 Softmax regression and SGD: implementation
- 3 Vectorization
- 4 Going deep: Activation functions

[Z, Section 4.4.7], Exercise 1

Numeric stability of `softmax` (see question 4(b) on TP2).
See [notebook](#).

The loss function in the case of binary classification.
See [notebook](#).

General remarks about the labs (TP)

- The lab exercises (TP) are an **essential** part of the course!

General remarks about the labs (TP)

- The lab exercises (TP) are an **essential** part of the course!
- Next week's TP 4 will be handed in and graded (25% of your grade).

General remarks about the labs (TP)

- The lab exercises (TP) are an **essential** part of the course!
- Next week's TP 4 will be handed in and graded (25% of your grade).
- In today's TP, you will make last week's solution more efficient using vectorization.

A (slow) implementation of softmax regression

See [notebook](#).

Today's lecture

- 1 Recap
- 2 Softmax regression and SGD: implementation
- 3 Vectorization**
- 4 Going deep: Activation functions

- Numpy arrays have a `shape` attribute.

- Numpy arrays have a `shape` attribute.
- Broadcasting happens when we perform an operation on two numpy arrays that do not have the same shape.

Broadcasting example

Example

(See [notebook.](#))

```
A = np.array([[1.],[2.],[3.],[4.]]) # shape: (4,1)
B = np.array([[5.,-5.,5.,-5.,5.]]) # shape: (1,5)
A + B
```

Images generated by [Google ImageFX](#), input prompts: "An unhappy mathematician, putting his head against a blackboard" and "A happy computer, which can perform any computation you ask of it".

Broadcasting example

Example

(See [notebook.](#))

```
A = np.array([[1.],[2.],[3.],[4.]]) # shape: (4,1)
B = np.array([[5.,-5.,5.,-5.,5.]]) # shape: (1,5)
A + B
```

Images generated by [Google ImageFX](#), input prompts: "An unhappy mathematician, putting his head against a blackboard" and "A happy computer, which can perform any computation you ask of it".

Broadcasting example

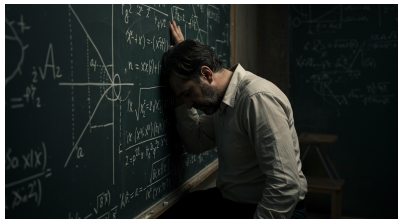
Example

(See [notebook](#).)

```
A = np.array([[1.],[2.],[3.],[4.]]) # shape: (4,1)
```

```
B = np.array([[5.,-5.,5.,-5.,5.]]) # shape: (1,5)
```

```
A + B
```



The Mathematician: Not allowed!!

Images generated by [Google ImageFX](#), input prompts: "An unhappy mathematician, putting his head against a blackboard" and "A happy computer, which can perform any computation you ask of it".

Broadcasting example

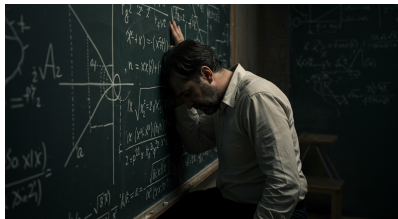
Example

(See [notebook](#).)

```
A = np.array([[1.],[2.],[3.],[4.]]) # shape: (4,1)
```

```
B = np.array([[5.,-5.,5.,-5.,5.]]) # shape: (1,5)
```

```
A + B
```



The Mathematician: Not allowed!!

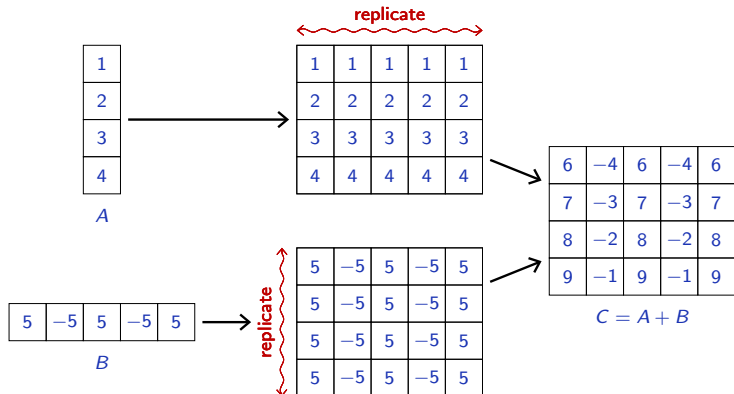


Numpy: This is fine.

Images generated by [Google ImageFX](#), input prompts: "An unhappy mathematician, putting his head against a blackboard" and "A happy computer, which can perform any computation you ask of it".

Broadcasting example visualized

```
A = torch.tensor([[1.], [2.], [3.], [4.]])  
B = torch.tensor([[5.], [-5.], [5.], [-5.], [5.]])  
C = A + B
```



Broadcasting visualized

$$\begin{array}{|c|c|c|} \hline (3,3) \\ \hline 1 & 2 & 3 \\ \hline 4 & 5 & 6 \\ \hline 7 & 8 & 9 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline (3,) \text{ or } (1,3) \\ \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline -1 & 0 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline (3,3) \\ \hline -1 & 0 & 3 \\ \hline -4 & 0 & 6 \\ \hline -7 & 0 & 9 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline (3,3) \\ \hline 1 & 2 & 3 \\ \hline 4 & 5 & 6 \\ \hline 7 & 8 & 9 \\ \hline \end{array} / \begin{array}{|c|c|c|} \hline (3,1) \\ \hline 3 & 3 & 3 \\ \hline 6 & 6 & 6 \\ \hline 9 & 9 & 9 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline (3,3) \\ \hline .3 & .7 & 1. \\ \hline .6 & .8 & 1. \\ \hline .8 & .9 & 1. \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline (3,) \text{ or } (1,3) \\ \hline 1 & 2 & 3 \\ \hline 1 & 2 & 3 \\ \hline 1 & 2 & 3 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline (3,1) \\ \hline 1 & 1 & 1 \\ \hline 2 & 2 & 2 \\ \hline 3 & 3 & 3 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline (3,3) \\ \hline 1 & 2 & 3 \\ \hline 2 & 4 & 6 \\ \hline 3 & 6 & 9 \\ \hline \end{array}$$

(From <https://towardsdatascience.com/broadcasting-in-numpy-58856f026d73>)

Broadcasting in practice: vectorized batch processing

- We represented our dataset as a 60000×784 matrix.

Broadcasting in practice: vectorized batch processing

- We represented our dataset as a 60000×784 matrix.
- We then processed each line separately, for example, for every batch, we repeat the following line 64 times:

```
probs = softmax(x @ W + b)
```


Broadcasting in practice: vectorized batch processing

- We represented our dataset as a 60000×784 matrix.
- We then processed each line separately, for example, for every batch, we repeat the following line 64 times:
`probs = softmax(x @ W + b)`
- Numpy (and Pytorch) work *faster* than Python loops.

Broadcasting in practice: vectorized batch processing

- We represented our dataset as a 60000×784 matrix.
- We then processed each line separately, for example, for every batch, we repeat the following line 64 times:

```
probs = softmax(x @ W + b)
```

- Numpy (and Pytorch) work *faster* than Python loops.
- We can **vectorize** the calculation `softmax(x @ W + b)`.

Batch matrix multiplication

- Instead of repeating $x @ W + b$ for every row x separately, we can take advantage of broadcasting, and do it at once for the entire batch:

`x.shape # (64, 784)`

`W.shape # (784, 10)`

`b.shape # (1, 10)`

`(x @ W).shape # (64, 10)`

`(x @ W + b).shape # (64, 10)`

Broadcasting details

- ❶ If one of the arrays has fewer dimensions than the other, then it is reshaped by adding dimension 1 on the left.

Broadcasting details

- ① If one of the arrays has fewer dimensions than the other, then it is reshaped by adding dimension 1 on the left.
- ② For every remaining dimension mismatch, *if one of the dimensions is equal to 1*, then it is expanded by replicating elements.

Broadcasting details

- ① If one of the arrays has fewer dimensions than the other, then it is reshaped by adding dimension 1 on the left.
- ② For every remaining dimension mismatch, *if one of the dimensions is equal to 1*, then it is expanded by replicating elements.
- ③ If there is a mismatch where both have a dimension not equal to 1, **fail**.

Broadcasting details

- 1 If one of the arrays has fewer dimensions than the other, then it is reshaped by adding dimension 1 on the left.
- 2 For every remaining dimension mismatch, *if one of the dimensions is equal to 1*, then it is expanded by replicating elements.
- 3 If there is a mismatch where both have a dimension not equal to 1, **fail**.

Reflection

When can we compute $A + B$?

Broadcasting details

- 1 If one of the arrays has fewer dimensions than the other, then it is reshaped by adding dimension 1 on the left.
- 2 For every remaining dimension mismatch, *if one of the dimensions is equal to 1*, then it is expanded by replicating elements.
- 3 If there is a mismatch where both have a dimension not equal to 1, **fail**.

Reflection

When can we compute $A + B$?

- `A.shape = (10)`, `B.shape = (5,1)`

Broadcasting details

- 1 If one of the arrays has fewer dimensions than the other, then it is reshaped by adding dimension 1 on the left.
- 2 For every remaining dimension mismatch, *if one of the dimensions is equal to 1*, then it is expanded by replicating elements.
- 3 If there is a mismatch where both have a dimension not equal to 1, **fail**.

Reflection

When can we compute $A + B$?

- $A.shape = (10), B.shape = (5, 1)$
- $A.shape = (10), B.shape = (1, 5)$

Broadcasting details

- 1 If one of the arrays has fewer dimensions than the other, then it is reshaped by adding dimension 1 on the left.
- 2 For every remaining dimension mismatch, *if one of the dimensions is equal to 1*, then it is expanded by replicating elements.
- 3 If there is a mismatch where both have a dimension not equal to 1, **fail**.

Reflection

When can we compute $A + B$?

- $A.shape = (10), B.shape = (5,1)$
- $A.shape = (10), B.shape = (1,5)$
- $A.shape = (1,10), B.shape = (1,10,1)$

Today's lecture

- 1 Recap
- 2 Softmax regression and SGD: implementation
- 3 Vectorization
- 4 Going deep: Activation functions

- The last additional ingredient: rectified linear unit (ReLU):

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0. \end{cases}$$

- The last additional ingredient: rectified linear unit (ReLU):

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0. \end{cases}$$

- Briefly: $\text{ReLU}(x) = \max(x, 0)$. ReLU is an *activation function*.

- The last additional ingredient: rectified linear unit (ReLU):

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0. \end{cases}$$

- Briefly: $\text{ReLU}(x) = \max(x, 0)$. ReLU is an *activation function*.
- Its derivative is very easy to compute:

$$\frac{\partial \text{ReLU}}{\partial x}(x_0) = \begin{cases} 1 & \text{if } x_0 \geq 0 \\ 0 & \text{if } x_0 < 0. \end{cases}$$

- The last additional ingredient: rectified linear unit (ReLU):

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0. \end{cases}$$

- Briefly: $\text{ReLU}(x) = \max(x, 0)$. ReLU is an *activation function*.
- Its derivative is very easy to compute:

$$\frac{\partial \text{ReLU}}{\partial x}(x_0) = \begin{cases} 1 & \text{if } x_0 \geq 0 \\ 0 & \text{if } x_0 < 0. \end{cases}$$

- Multi-layer perceptron = stack of several linear layers, with ReLU in between each layer.

Numeric stability trick

- We saw that, for numeric stability, instead of computing

$$\text{softmax}(\mathbf{o})_j = \frac{o_j}{P} \text{ where } P = \sum_j \exp(o_j)$$

we compute, for a number $M \in \mathbb{R}$,

$$\text{softmax}'(\mathbf{o})_j = \frac{o_j - M}{P} \text{ where } P = \sum_j \exp(o_j - M)$$

Numeric stability trick

- We saw that, for numeric stability, instead of computing

$$\text{softmax}(\mathbf{o})_j = \frac{o_j}{P} \text{ where } P = \sum_j \exp(o_j)$$

we compute, for a number $M \in \mathbb{R}$,

$$\text{softmax}'(\mathbf{o})_j = \frac{o_j - M}{P} \text{ where } P = \sum_j \exp(o_j - M)$$

Reflection

Why does this give the same answer?

Computing the gradient of softmax by hand

- We compute the logits $\mathbf{o} = \mathbf{x}^T \cdot \mathbf{\Omega} + \beta$.

Computing the gradient of softmax by hand

- We compute the logits $\mathbf{o} = \mathbf{x}^T \cdot \mathbf{\Omega} + \beta$.
- We compute the probability vector $\mathbf{p} = \text{softmax}(\mathbf{o})$.

Computing the gradient of softmax by hand

- We compute the logits $\mathbf{o} = \mathbf{x}^T \cdot \mathbf{\Omega} + \beta$.
- We compute the probability vector $\mathbf{p} = \text{softmax}(\mathbf{o})$.
- If the correct class is c , the loss on \mathbf{x} is $\ell \stackrel{\text{def}}{=} -\log p_c$.

Computing the gradient of softmax by hand

- We compute the logits $\mathbf{o} = \mathbf{x}^T \cdot \mathbf{\Omega} + \beta$.
- We compute the probability vector $\mathbf{p} = \text{softmax}(\mathbf{o})$.
- If the correct class is c , the loss on \mathbf{x} is $\ell \stackrel{\text{def}}{=} -\log p_c$.
- We saw last week the **gradient** of this new loss function:

$$\mathbf{p}' \stackrel{\text{def}}{=} \frac{\partial \ell}{\partial \mathbf{o}} = \begin{cases} p_j & \text{if } j \neq c \\ p_j - 1 & \text{if } j = c. \end{cases}$$

Computing the gradient of softmax by hand

- We compute the logits $\mathbf{o} = \mathbf{x}^T \cdot \mathbf{\Omega} + \beta$.
- We compute the probability vector $\mathbf{p} = \text{softmax}(\mathbf{o})$.
- If the correct class is c , the loss on \mathbf{x} is $\ell \stackrel{\text{def}}{=} -\log p_c$.
- We saw last week the **gradient** of this new loss function:

$$\mathbf{p}' \stackrel{\text{def}}{=} \frac{\partial \ell}{\partial \mathbf{o}_j} = \begin{cases} p_j & \text{if } j \neq c \\ p_j - 1 & \text{if } j = c. \end{cases}$$

Reflection

Why is this the gradient?

TODO for next week

- Finish TP3.

TODO for next week

- Finish TP3.
- Read [P, Chapter 3].

TODO for next week

- Finish TP3.
- Read [P, Chapter 3].
- Review the reflection questions on the last two slides above.

[P] Prince, S., **Understanding Deep Learning**, <https://udlbook.com>

[Z] Zhang, A. et al., **Dive into Deep Learning**, <https://www.d2l.ai>

Image sources:

- <https://en.wikipedia.org/wiki/Omega>
- [https://en.wikipedia.org/wiki/Nevel_\(instrument\)](https://en.wikipedia.org/wiki/Nevel_(instrument))